


Comparative Analysis of ANN and CNN Models

Exploring the performance of Artificial Neural Networks (ANN) and Convolutional Neural Networks (CNN) for regression and classification tasks using PyTorch and Keras.

 by Sadia Shakoor



Datasets

California Housing

Regression, 20,640 records, 8 features, median house value target.

Customer Churn

Classification, 7,043 records, 20 features, binary churn target.

CIFAR-10

Classification, 60,000 images, 32x32 RGB, 10 classes.

PyTorch Model Architectures

1

PyTorch ANN (Regression)

32-16-8-1 layers, ReLU activation.

2

PyTorch ANN (Classification)

16-8-1 layers, ReLU activation, sigmoid output.

CNN Model Architectures

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 8)	224
max_pooling2d (MaxPooling2D)	(None, 16, 16, 8)	0
conv2d_1 (Conv2D)	(None, 16, 16, 16)	1,168
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 16)	0
conv2d_2 (Conv2D)	(None, 8, 8, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 32)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 128)	65,664
dense_1 (Dense)	(None, 10)	1,290

Total params: 72,986 (285.10 KB)
Trainable params: 72,986 (285.10 KB)
Non-trainable params: 0 (0.00 B)

Model 1: Baseline CNN

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	73,856
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_2 (Dense)	(None, 128)	262,272
dense_3 (Dense)	(None, 10)	1,290

Total params: 356,810 (1.36 MB)
Trainable params: 356,810 (1.36 MB)
Non-trainable params: 0 (0.00 B)

Model 2: Increased Filter Size

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_12 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_13 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_13 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_14 (Conv2D)	(None, 8, 8, 128)	73,856
max_pooling2d_14 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_4 (Flatten)	(None, 2048)	0
dense_10 (Dense)	(None, 128)	262,272
dense_11 (Dense)	(None, 64)	8,256
dense_12 (Dense)	(None, 10)	650

Total params: 364,420 (1.39 MB)
Trainable params: 364,420 (1.39 MB)
Non-trainable params: 0 (0.00 B)

Model 3: Additional Dense Layer

Training Configurations

Regression

SGD optimizer, learning rates 0.01, 0.1, epochs 30, 50, batch sizes $\text{len}(\text{X_train})/10$, $\text{len}(\text{X_train})/5$ and $\text{len}(\text{X_train})$.

Classification

Adam optimizer, learning rates 0.01, 0.1, epochs 30, 100, batch sizes $\text{len}(\text{X_train})/10$, $\text{len}(\text{X_train})/5$ and $\text{len}(\text{X_train})$.

CNN

Adam optimizer, learning rate 0.001, epochs 10, 15, 20, batch size 32 and 64.

Performance Metrics

1.31

MSE

Regression, epoch 30, learning rate 0.01.

83%

Accuracy

Classification, epoch 100, learning rate 0.01.

0.90

MAE

Regression, epoch 30, learning rate 0.01.

74%

Accuracy

CNN Model 2, epoch 20, learning rate 0.001.

Visualizations



Learning Curves


Training and validation loss, accuracy over epochs.



Confusion Matrices

Class-wise performance for classification tasks.

Learning Curves for Regression Models



Model 1

Optimizer =SGD

Learning Rate =0.01


Loss Fun = MSELoss

Epoch=30

Batch Size=int(len(X_train)/10)

Mean Squared Error: 1.3105

Mean Absolute Error:0.90097



Model 2

Optimizer =SGD

Learning Rate =0.01

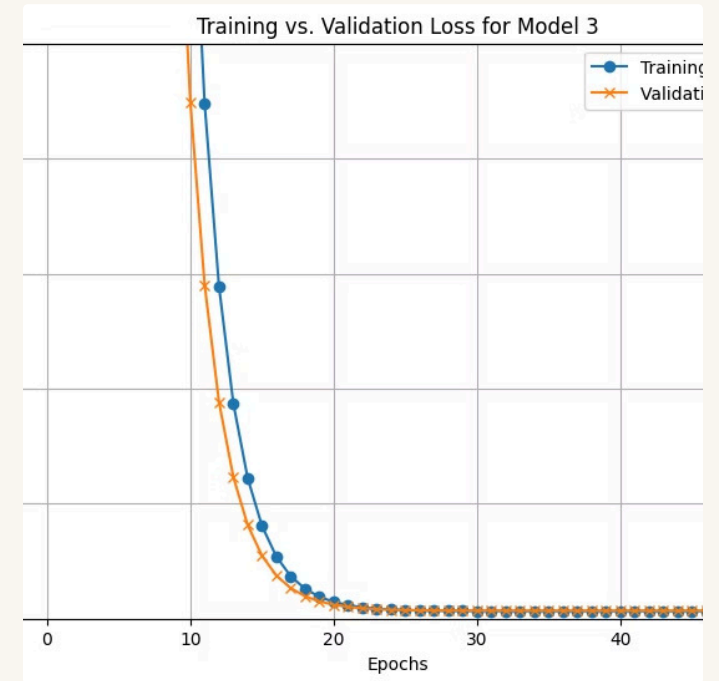
Loss Fun = MSELoss

Epoch=100

Batch Size=int(len(X_train)/5)

Mean Squared Error: 1.3218

Mean Absolute Error: 0.9261



Model 3

Optimizer =SGD

Learning Rate =0.01

Loss Fun = MSELoss


Epoch=100

Batch Size=int(len(X_train))

Mean Squared Error: 1.3218

Mean Absolute Error: 0.92607

Confusion Matrices for Classification Model



Model 1

Optimizer = Adam


Learning Rate = 0.01

Loss Fun = BCELoss

Epoch= 30

Batch Size=int(len(X_train)/10)

Accuracy: 0.7395



Model 2

Optimizer = Adam


Learning Rate = 0.01

Loss Fun = BCELoss

Epoch= 100

Batch Size=int(len(X_train)/5)

Accuracy: 0.8325



Model 3

Optimizer = Adam

Learning Rate = 0.1


Loss Fun = BCELoss

Epoch= 100

Batch Size=int(len(X_train))

Accuracy: 0.74

Confusion Matrices for CNN Models



Model 1

Optimizer = Adam


Learning Rate =0.01

Loss Fun = BCELoss

Epoch=10

Batch Size=32

Accuracy= 0.68 10000



Model 2

Optimizer = Adam


Learning Rate =0.01

Loss Fun = BCELoss

Epoch= 20

Batch Size= 64

Accuracy= 0.74



Model 3

Optimizer = Adam

Learning Rate =0.001

Loss Fun = BCELoss

Epoch= 15

Batch Size= 32

Accuracy= 0.73

Comparative Table

Model	Dataset / Task	Key Hyperparams	Final Metric	Training Time
PyTorch ANN (Reg)	California Housing	LR=0.01, Epoch=30	MSE=1.31, MAE=0.90	~1 min
PyTorch ANN (Class)	Customer Churn	LR=0.01, Epoch=100	Accuracy=83%, Recall=52%	~1 min
Keras CNN (Model 2)	CIFAR-10	LR=0.001, Epoch=20	Accuracy=74%	~1 min (GPU)

Key Takeaways

1

Strengths

PyTorch ANN: Simple architectures, decent results for regression and binary classification.

2

Strengths

Keras CNN: Improved accuracy with deeper architectures and dropout regularization.

3

Weaknesses

PyTorch ANN: Limited recall for imbalanced datasets.

4

Weaknesses

Keras CNN: Performance plateaued with increased complexity.